



SAPIENZA
UNIVERSITÀ DI ROMA

Homework 2

MACHINE LEARNING

Students:

Flavio Maiorana

2051396

Academic Year 2023/2024

1 Introduction

First of all, it could be useful to gain some insight on how the dataset is made. The available data is already split into training and testing set. We will consider the test split only in the evaluation phase (to prevent data leakage).

[Dataset for Training]

```
N Examples: 6369
N Classes: 5
Classes: [0 1 2 3 4]
- Class 0: 1000 (15.701051970482022)
- Class 1: 1500 (23.551577955723033)
- Class 2: 1500 (23.551577955723033)
- Class 3: 2000 (31.402103940964043)
- Class 4: 369 (5.7936881771078665)
```

Some comments: the dataset is highly imbalanced. That needs to be addressed accordingly. The main two problems to address are:

- Model architecture
- Training process
 - Class imbalance mitigation
 - Optimization choice

Before everything else, it is worth to mention the performance metrics. Since the dataset is highly imbalanced, accuracy might not be the best choice to measure performance. Rather, the weighted mean of f1-scores (and also recall and precision) could be a better choice.

2 Training process

There are some important issues to address, which are to some extent independent from the model we use to classify. First of all, to mitigate class imbalance we could apply:

- Data augmentation (useful in any case to increase generalization)
- Different sampling techniques
 - Upsampling/Downsampling
 - Weighted resampling

2.1 Sampling techniques

Different sampling of data in order to rebalance the classes. These techniques use also data augmentation techniques, in order to make the network as generalizing as possible. Data augmentation is used obviously also without resampling techniques.

2.2 Data augmentation

This technique is widely used with images. It exploits the invariance property of images to enhance generalization properties of neural networks. In our case it is particularly useful, because it allows us to "reuse" the same sample multiple times and make it look new to the network by augmenting its features. The transformation I used are these three from the torchvision package, plus a random color jitter implemented separately. I also tried RandomHorizontalFlipping, but I noticed actually worse conditions, maybe because the images are sensitive to the horizontal orientation (for turning left and right), so randomly flipping them would mean introducing noise to the labels.

```
transforms.RandomInvert(),  
transforms.RandomErasing(),  
transforms.RandomRotation(degrees=5),
```

2.2.1 Weighted random resampling

In order to mitigate class imbalance, one could try to rebalance the dataset by using weighted resampling techniques. This can be achieved by `torch.data.WeightedRandomSampler`. The weights can be chosen according to class distribution, but I achieved better performance by tuning them 'by hand'. More in particular, the weights have to be chosen such that every batch has an equal distribution among all classes.

```
[0.8,0.6,0.6,0.5,1.7]
```

For example, after some trials, the above weights vector proved itself to perform better than without weighted resampling.

Naturally, class 3 has to be sampled with the least probability, and class 4 with the highest probability. Besides, the data will be resampled with replacement. I was able to reach a satisfactory performance with some manual tuning of the class weights, but the problem of this approach is that the `WeightedRandomSampler` is not perfectly reproducible, which means that performance from one trained model to the other. Nevertheless, I was able to complete the entire lap only with this method.

2.2.2 Over/undersampling

Another useful and potentially effective technique to make up for class imbalance is over and undersampling. To that end, I used the python package `imblearn`, that

offers out-of-the-box methods to under/over sample a dataset. I used the SMOTEENN method, which first oversamples all but the minority class by 'generating' new samples through interpolation between two existing samples. After that, ENN is applied, which polishes the newly generated samples, removing the nearest one to each other. After over and under sampling the new dataset will look like this. We lost some samples, but now the dataset is more balanced and well distributed.

```
N Examples: 3399
N Classes: 5
Classes: [0 1 2 3 4]
- Class 0: 515 (15.151515151515152)
- Class 1: 628 (18.476022359517504)
- Class 2: 658 (19.358634892615477)
- Class 3: 903 (26.5666372462489)
- Class 4: 695 (20.44719035010297)
```



Figure 1: Different methods compared

Anytime an image is sampled, these transforms are applied, resulting in a completely new image.

2.2.3 Comparing them all

In the end, from the point of view of the f1-score the performances are all similar, maybe smoteenn has a slightly better performance. But in the end, on the track the method using nothing always behaves badly, the others do better. The video uses weighted random resampling.

2.3 Optimizer

I tested two different optimizers: SGD and Adam. The second one had generally a better performance, resulting in a smoother optimization process. The best learning rate was 0.0001. A bigger lr resulted in a 'ripple' effect on the learning process, due to the 'jumps'.

2.3.1 Early stopping

One could also implement a stopping mechanism. That is helpful to stop the loss from sinking while having a performance measurement on a validation set, which means we are overfitting on the training data. The validation split needs to be totally independent from training data, namely no data augmentation on it. Regarding early stopping, it is important to adequately choose the 'patience' parameter. It indicates the number of epochs after which, if the validation score is getting worse, training will stop. In general, early stopping can be applied both to validation score and loss. In my case, a patience lower than 3 resulted in too early stopping, and a patience above 10 resulted in sometimes overfitting.

3 Model Architecture

The first and foremost choice regarding cnn model architecture is surely how deep it has to be and how big the kernel need to be. Since the images are not that big, a reasonably small cnn would be enough, since a too deep one would either overfit on training data, or waste computational resources without gaining any improved performance. Although, a too small model would not be able to adequately catch the features of the image. Also, kernels size is very important. A too big kernel would compromise computational efficiency, but the size has to be adapted to the type of features the convolution needs to capture. To corroborate these concepts we will compare some CNNs with a different number of convolutional layers.

Each CNN has similar structure: N convolutional layers, each followed by a batch normalization layer, then an activation function and at last a max pooling layer. After the convolution layers, responsible for extracting features from the image, the network has some fully connected layers (with the same activation function as the convolutional

layers) at the end, followed by a final dropout layer, that have as final output a vector of five elements, expressing the score predicted for each class.

In the following part, each list entry indicates the dimension for the convolutional layer (from 1 to N). The kernel size 0 indicates the number of channels of the input image. This training results are all captured after 20 epochs.

- $N = 4$

```
'channels': [3,5,15,30,30],  
'kernels': [7,5,5,5],  
'strides': [1,1,1,1],  
'pool_kernels': [2,2,2,2],  
'pool_strides': [2,2,2,2],  
'fc_dims': [120,5],  
'dropout': 0.2  
Training time: 169.67
```

- $N = 6$

```
'channels': [3,10,16,32,64,64,64],  
'kernels': [5,5,5,3,3,3],  
'strides': [1,1,1,1,1,1],  
'pool_kernels': [2,2,2,2,2,2],  
'pool_strides': [2,1,2,1,2,1],  
'fc_dims': [576,64,5],  
'dropout': 0.2  
Training time: 400.342
```




Figure 2: Deeper network takes more to train and overfits training data

3.1 Dropout layer




Figure 3: Dropout

This layer is inserted at the end of the fully connected layers. It is particularly important to avoid overfitting. More specifically, it avoids that the model overfits too much on training data by inserting some perturbations in the model, namely by dropping randomly some connections from one layer to the other. As we can see, with the same exact architecture, just by eliminating the dropout layer, the learning becomes much different, and so also the result is different, as we can see from the two confusion matrices.




Figure 4: Dropout confusion matrix




Figure 5: No Dropout confusion matrix

The inference tends more towards class 3, which means overfitting.

4 Final Considerations

In the end, after different trials I was able to achieve the completion of a full lap through manual tuning of class weights. The difficulty in this task was that a relatively high score, whatever the metrics, not always reflected in a high reward on the real track. Moreover, I noticed that sometimes the car did better on the track when the model never predicted class 0 or 4, because the risk is that a false positive of one of these two would stop the car for the entire episode. So, in general precision needs to be high especially for these two classes, and that can be regulated better with weighted resampling.