

# Deep Learning

03 - Model Training. Loss function



SAPIENZA  
UNIVERSITÀ DI ROMA

Fabrizio Silvestri

# Recap





# Recapping

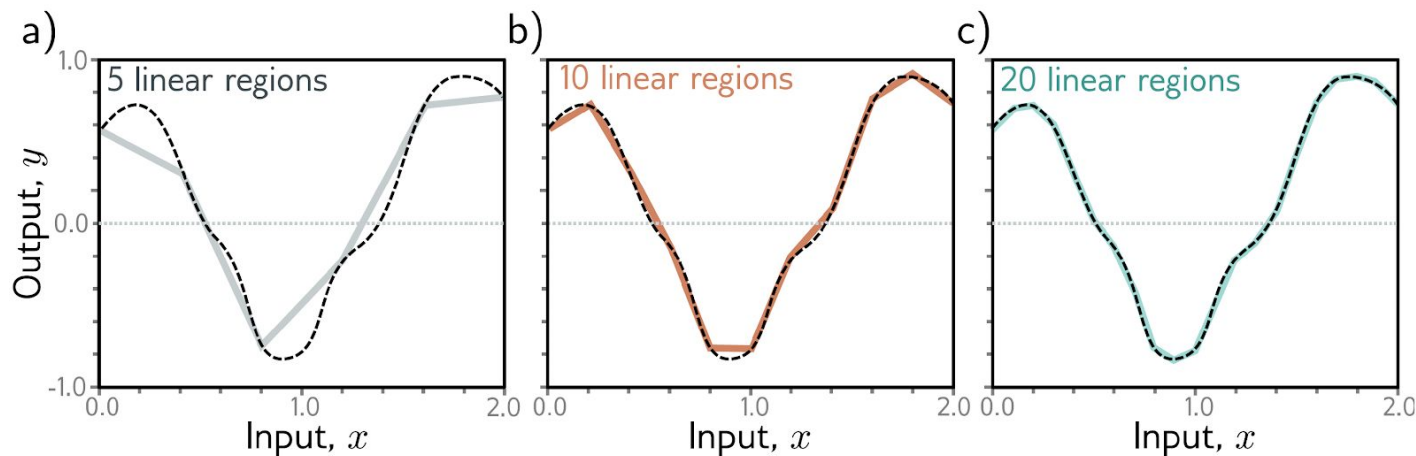
- We have described linear regression, shallow neural networks, and deep neural networks.
- Each represents a family of functions that map input to output, where the particular member of the family is determined by the model parameters  $\phi$ .
- When we train these models, we seek the parameters that produce the **best possible mapping** from input to output for the task we are considering.
- What does “best possible” mapping mean?



# Neural Networks



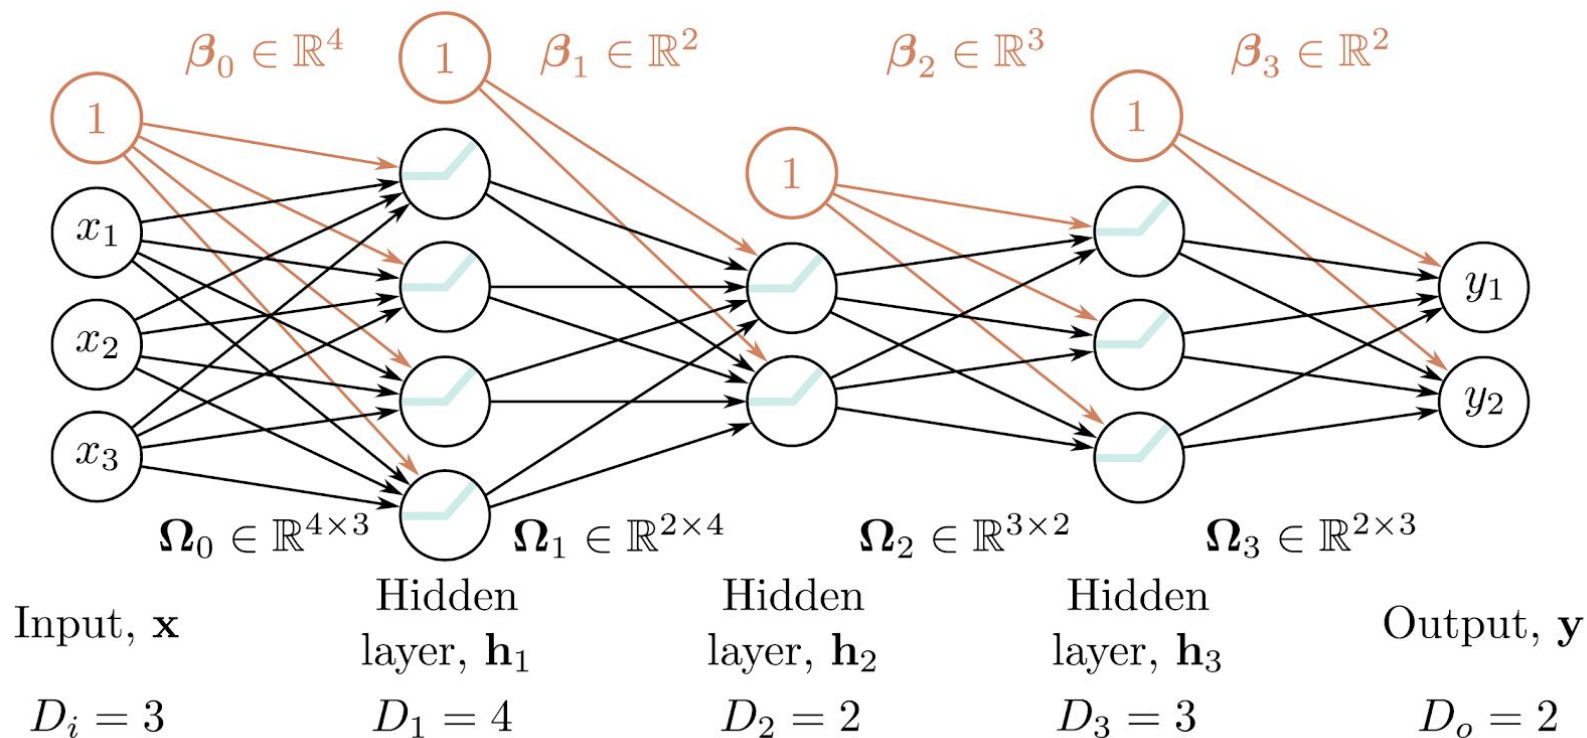
# Universal Approximation Theorem




**Figure 3.5** Approximation of a 1D function (dashed line) by a piecewise linear model. a–c) As the number of regions increases, the model becomes closer and closer to the continuous function. A neural network with a scalar input creates one extra linear region per hidden unit. The universal approximation theorem proves that, with enough hidden units, there exists a shallow neural network can describe any given continuous function defined on a compact subset of  $\mathbb{R}^D$  to arbitrary precision.




# More Hidden Layers






# More Hidden Layers



# Decision Boundary



- 2 hidden layer
  - Combinations of convex regions






# Why Deep Learning

- Features that are engineered by hand are difficult to maintain, error prone, and not necessarily the best for the task
- Can we learn feature directly from data?

**Low Level Features**



Lines & Edges

**Mid Level Features**



Eyes & Nose & Ears

**High Level Features**



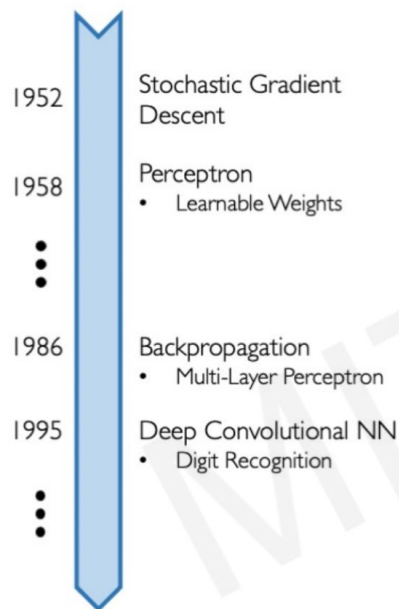
Facial Structure



# Why Now?

- ANNs are not new (The Perceptron is almost 65 y.o.)

Neural Networks date back decades, so why the resurgence?



## 1. Big Data

- Larger Datasets
- Easier Collection & Storage

IMAGENET



## 2. Hardware

- Graphics Processing Units (GPUs)
- Massively Parallelizable



## 3. Software

- Improved Techniques
- New Models
- Toolboxes



# Loss Functions

(Chapter 5)



# Requirements

- That definition requires a training dataset  $\{x_i, y_i\}$  of input/output pairs.
- A loss function, or cost function,  $L[\phi]$  returns a single number that describes the mismatch between the model predictions  $f[x_i, \phi]$  and their corresponding ground-truth outputs  $y_i$ .
- During training, we seek parameter values  $\tilde{\phi}$  that minimize the loss and hence map the training inputs to the outputs as closely as possible.

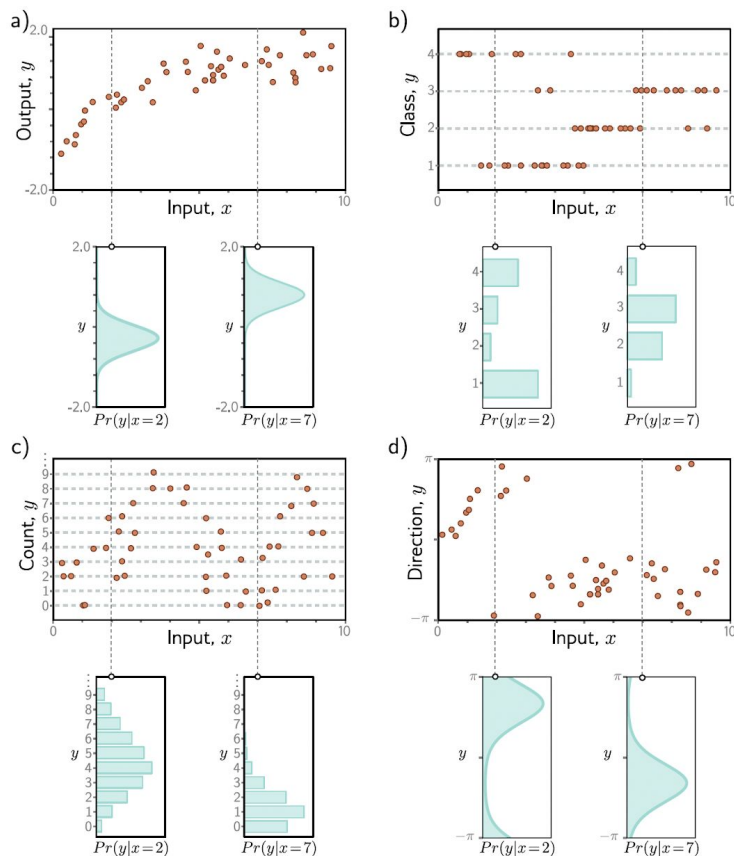
$$\tilde{\phi} = \arg \min_{\phi} L[\phi]$$



# Maximum Likelihood



# Models as Conditional Probability Distributions



**Figure 5.1** Predicting distributions over outputs. a) Regression task, where the goal is to predict a real-valued output  $y$  from the input  $x$  based on training data  $\{x_i, y_i\}$  (orange points). For each input value  $x$ , the machine learning model predicts a distribution  $Pr(y|x)$  over the output  $y \in \mathbb{R}$  (cyan curves show distributions for  $x=2.0$  and  $x=7.0$ ). The loss function aims to maximize the probability of the observed training outputs  $y_i$  under the distribution predicted from the corresponding inputs  $x_i$ . b) To predict discrete classes  $y \in \{1, 2, 3, 4\}$  in a classification task, we use a discrete probability distribution, so the model predicts a different histogram over the four possible values of  $y_i$  for each value of  $x_i$ . c) To predict counts  $y \in \{0, 1, 2, \dots\}$  and d) direction  $y \in (-\pi, \pi]$ , we use distributions defined over positive integers and circular domains, respectively.



# Computing a Distribution over Outputs

- How do we do that?
- Simple
  - $y=f[x, \phi] \Rightarrow \Pr\{y|x; \phi\}$

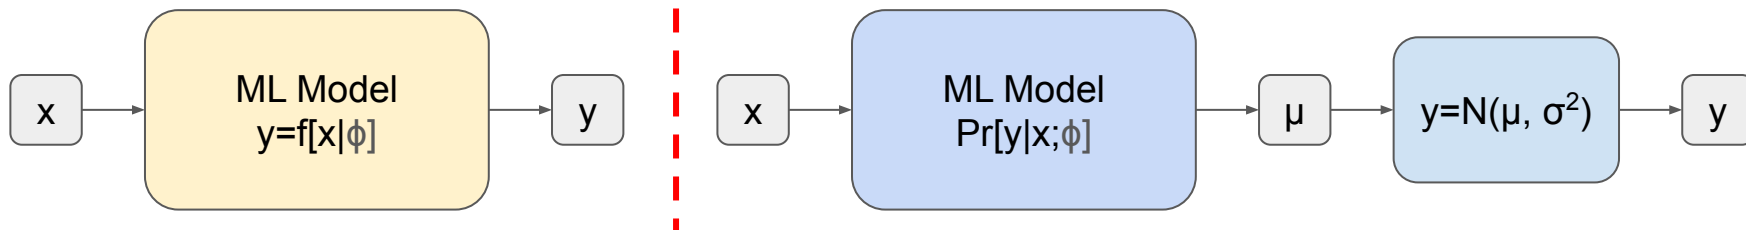


Image courtesy of: <https://www.polygon.com/22947332/spider-man-no-way-home-pointing-meme>



# Example


- Suppose the prediction domain is the set of real numbers, i.e.,  $y \in \mathbb{R}$ .
- Here, we might choose the univariate normal distribution, which is defined on  $\mathbb{R}$ .
  - This distribution is defined by the mean  $\mu$  and variance  $\sigma^2$ , so  $\theta = \{\mu, \sigma^2\}$ .
- The machine learning model might predict the mean  $\mu$ , and the variance  $\sigma^2$  could be treated as an unknown constant.





# MLE Criterion

- The model now computes different distribution parameters  $\theta_i = f[x_i, \phi]$  for each training input  $x_i$ .
  - Each observed training output  $y_i$  should have high probability under its corresponding distribution  $\Pr\{y_i|\theta_i\}$ .
- Hence, we choose the model parameters  $\phi$  so that they maximize the combined probability across all  $i$  training examples



$$\begin{aligned}\hat{\phi} &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I \Pr(y_i|x_i) \right] \\ &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I \Pr(y_i|\theta_i) \right] \\ &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I \Pr(y_i|f[x_i, \phi]) \right]\end{aligned}$$



# MLE Criterion: Assumptions

- Observations  $(\mathbf{x}_i, y_i)$  are:
  - Independent
  - Identically distributed
- Observations i.i.d.




$$Pr(\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_I | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_I) = \prod_{i=1}^I Pr(\mathbf{y}_i | \mathbf{x}_i)$$

$$\begin{aligned}\hat{\phi} &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I Pr(\mathbf{y}_i | \mathbf{x}_i) \right] \\ &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I Pr(\mathbf{y}_i | \boldsymbol{\theta}_i) \right] \\ &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I Pr(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right]\end{aligned}$$



# Log Likelihood

$$\begin{aligned}\hat{\phi} &= \operatorname{argmax}_{\phi} \left[ \prod_{i=1}^I \operatorname{Pr}(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \\ &= \operatorname{argmax}_{\phi} \left[ \log \left[ \prod_{i=1}^I \operatorname{Pr}(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \right] \\ &= \operatorname{argmax}_{\phi} \left[ \sum_{i=1}^I \log \left[ \operatorname{Pr}(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \right]\end{aligned}$$



# Minimizing Negative Log Likelihood (NLL)

- We note that, by convention, model fitting problems are framed in terms of minimizing a loss.
- To convert the maximum log-likelihood criterion to a minimization problem, we multiply by minus one, which gives us the negative log-likelihood criterion:

$$\begin{aligned}\hat{\phi} &= \operatorname{argmin}_{\phi} \left[ \ominus \sum_{i=1}^I \log \left[ \operatorname{Pr}(\mathbf{y}_i | \mathbf{f}[\mathbf{x}_i, \phi]) \right] \right] \\ &= \operatorname{argmin}_{\phi} \left[ L[\phi] \right],\end{aligned}$$



# Inference

- The network no longer directly predicts the outputs  $y$  but instead determines a probability distribution over  $y$ .
- When we perform inference, we often want a point estimate rather than a distribution, so we return the maximum of the distribution:

$$\hat{\mathbf{y}} = \underset{\mathbf{y}}{\operatorname{argmax}} \left[ \operatorname{Pr}(\mathbf{y} | \mathbf{f}[\mathbf{x}, \hat{\phi}]) \right]$$

- It is usually possible to find an expression for this in terms of the distribution parameters  $\theta$  predicted by the model.
- For example, in the univariate normal distribution, the maximum occurs at the mean  $\mu$



# Recipe for a Loss

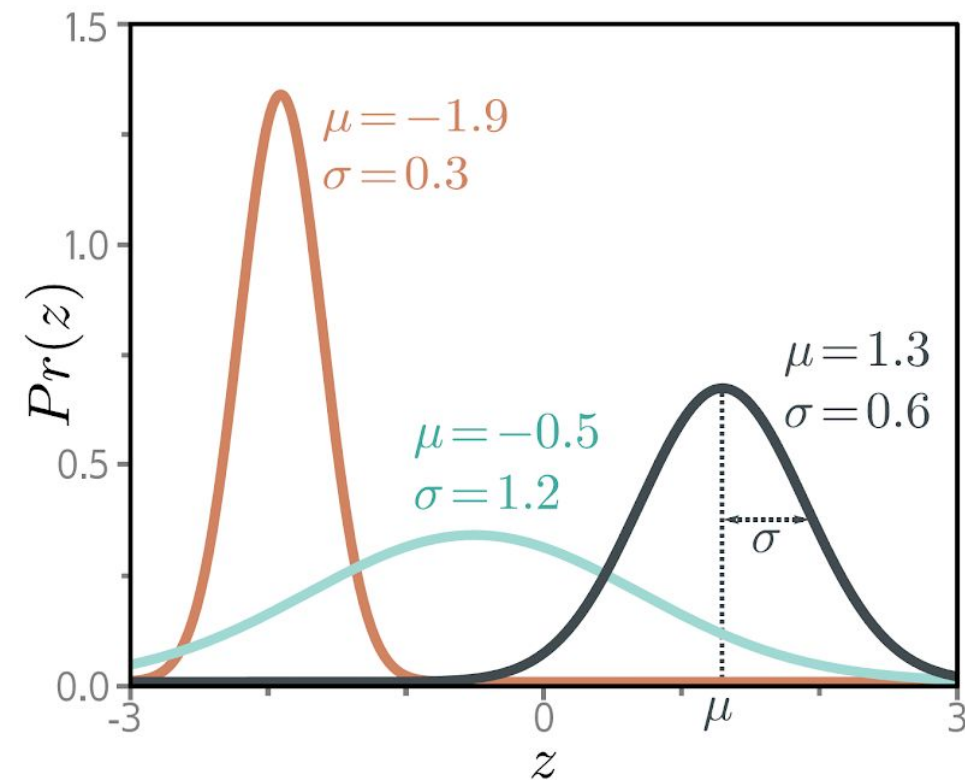
1. Choose a suitable probability distribution  $\Pr(y|\theta)$  defined over the domain of the predictions  $y$  with distribution parameters  $\theta$ .
2. Set the machine learning model  $f[x, \phi]$  to predict one or more of these parameters, so  $\theta = f[x, \phi]$  and  $\Pr(y|\theta) = \Pr(y|f[x, \phi])$ .
3. To train the model, find the network parameters  $\hat{\phi}$  that minimize the negative log-likelihood loss function over the training dataset pairs  $\{x_i, y_i\}$ :

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} [L[\phi]] = \underset{\phi}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log [Pr(y_i | f[x_i, \phi])] \right]$$

4. To perform inference for a new test example  $x$ , return either the full distribution  $\Pr(y|f[x, \hat{\phi}])$  or the maximum of this distribution.



# Example 1: Univariate Regression



**Figure 5.3** The univariate normal distribution (also known as the Gaussian distribution) is defined on the real line  $z \in \mathbb{R}$  and has parameters  $\mu$  and  $\sigma^2$ . The mean  $\mu$  determines the position of the peak. The positive root of the variance  $\sigma^2$  (the standard deviation) determines the width of the distribution. Since the total probability density sums to one, the peak becomes higher as the variance decreases and the distribution becomes narrower.



# Example 1: Univariate Regression

- Goal: predict a single scalar output  $y \in \mathbb{R}$  from input  $x$  using a model  $f[x, \phi]$  with parameters  $\phi$ .
- Following the recipe, we choose a probability distribution over the output domain  $y$

$$Pr(y|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y - \mu)^2}{2\sigma^2} \right]$$

pdf  
Probability Density Function

- Second, we set the machine learning model  $f[x, \phi]$  to compute one or more of the parameters of this distribution. Here, we just compute the mean so  $\mu = f[x, \phi]$ :

$$Pr(y|f[x, \phi], \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y - f[x, \phi])^2}{2\sigma^2} \right]$$





# Example 1: Univariate Regression

- We aim to find the parameters  $\phi$  that make the training data  $\{\mathbf{x}_i, y_i\}$  most probable under this distribution.
- To accomplish this, we choose a loss function  $L[\phi]$  based on the negative log-likelihood:

$$\begin{aligned} L[\phi] &= - \sum_{i=1}^I \log [Pr(y_i | f[\mathbf{x}_i, \phi], \sigma^2)] \\ &= - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \end{aligned}$$

NLLLoss

- When we train the model, we seek parameters  $\hat{\phi}$  that minimize this loss



# Example 1: Univariate Regression

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \right]$$



# Example 1: Univariate Regression

$$\begin{aligned}\hat{\phi} &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \right] \\ &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \right] - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right]\end{aligned}$$



# Example 1: Univariate Regression

$$\begin{aligned}\hat{\phi} &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \right] \\ &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \right] - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \\ &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right]\end{aligned}$$



# Example 1: Univariate Regression

$$\begin{aligned}\hat{\phi} &= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \right] \\&= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \right] - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \\&= \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I - \frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \\&= \operatorname{argmin}_{\phi} \left[ \sum_{i=1}^I (y_i - f[\mathbf{x}_i, \phi])^2 \right]\end{aligned}$$

MSELoss



# Colab time: Least Squares Loss

<https://colab.research.google.com/drive/1gy3KiCiKl4Zol2eKgH9XlzZif-S0hUPB>



# What about the Variance?

- Treating  $\sigma^2$  as a parameter of the model and minimizing the loss with respect to both the model parameters  $\phi$  and the distribution variance  $\sigma^2$ :

$$\hat{\phi} = \underset{\phi, \sigma^2}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left[ -\frac{(y_i - f[\mathbf{x}_i, \phi])^2}{2\sigma^2} \right] \right] \right]$$

- At inference time, the model predicts the mean  $\mu = f[\mathbf{x}, \hat{\phi}]$  from the input, and we **learned the variance  $\sigma^2$  during the training process.**
- The former is the best prediction. The latter tells us about the *uncertainty of the prediction.*



# Heteroscedastic Regression

- When the uncertainty of the model varies as a function of the input data, we refer to this as heteroscedastic.
- Train  $\mathbf{f}[\mathbf{x}, \boldsymbol{\phi}] = [f_1[\mathbf{x}, \boldsymbol{\phi}], f_2[\mathbf{x}, \boldsymbol{\phi}]]$ 
  - $f_2[\mathbf{x}, \boldsymbol{\phi}]$  can be negative  $\rightarrow$  square it

$$\begin{aligned}\mu &= f_1[\mathbf{x}, \boldsymbol{\phi}] \\ \sigma^2 &= f_2[\mathbf{x}, \boldsymbol{\phi}]^2\end{aligned}$$

- which results in the loss function:


$$\hat{\boldsymbol{\phi}} = \underset{\boldsymbol{\phi}}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log \left[ \frac{1}{\sqrt{2\pi f_2[\mathbf{x}_i, \boldsymbol{\phi}]^2}} - \frac{(y_i - f_1[\mathbf{x}_i, \boldsymbol{\phi}])^2}{2 f_2[\mathbf{x}_i, \boldsymbol{\phi}]^2} \right] \right]$$






# Heteroscedastic vs. Homoscedastic Regression


a)




b)



c)



d)



# Example 2: Binary Classification

- In binary classification, the goal is to assign the data  $x$  to one of two discrete classes  $y \in \{0, 1\}$ . In this context, we refer to  $y$  as a label.
- Examples of binary classification include:
  - predicting whether a restaurant review is positive ( $y = 1$ ) or negative ( $y = 0$ ) from text data  $x$
  - predicting whether a tumor is present ( $y = 1$ ) or absent ( $y = 0$ ) from an MRI scan  $x$



# Reminder: Recipe for a Loss

1. Choose a suitable probability distribution  $\Pr(y|\theta)$  defined over the domain of the predictions  $y$  with distribution parameters  $\theta$ .
2. Set the machine learning model  $f[x, \phi]$  to predict one or more of these parameters, so  $\theta = f[x, \phi]$  and  $\Pr(y|\theta) = \Pr(y|f[x, \phi])$ .
3. To train the model, find the network parameters  $\hat{\phi}$  that minimize the negative log-likelihood loss function over the training dataset pairs  $\{x_i, y_i\}$ :

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} [L[\phi]] = \underset{\phi}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log [Pr(y_i | f[x_i, \phi])] \right]$$

4. To perform inference for a new test example  $x$ , return either the full distribution  $\Pr(y|f[x, \hat{\phi}])$  or the maximum of this distribution.




# Example 2: Binary Classification

- Let's build the loss function, then...
- First, we choose a probability distribution over the output space  $y \in \{0, 1\}$ . A suitable choice is the Bernoulli distribution, which is defined on the domain  $\{0, 1\}$ . This has a single parameter  $\lambda \in [0, 1]$  that represents the probability that  $y$  takes the value one

$$Pr(y|\lambda) = \begin{cases} 1 - \lambda & y = 0 \\ \lambda & y = 1 \end{cases}$$

- which can equivalently be written as:


$$Pr(y|\lambda) = (1 - \lambda)^{1-y} \cdot \lambda^y$$



## Example 2: Binary Classification

- Second, we set the machine learning model  $f[\mathbf{x}, \boldsymbol{\phi}]$  to predict the single distribution parameter  $\lambda$ .
- However,  $\lambda$  can only take values in the range  $[0, 1]$ , and we cannot guarantee that the network output will lie in this range.
- Consequently, we pass the network output through a function that maps the real numbers  $\mathbb{R}$  to  $[0, 1]$ . A suitable function is the logistic sigmoid

$$\text{sig}[z] = \frac{1}{1 + \exp[-z]}$$



- Hence, we predict the distribution parameter as  $\lambda = \text{sig}[f[\mathbf{x}, \boldsymbol{\phi}]]$



## Example 2: Binary Classification

- The likelihood is, therefore:

$$Pr(y|\mathbf{x}) = (1 - \text{sig}[f[\mathbf{x}, \phi]])^{1-y} \cdot \text{sig}[f[\mathbf{x}, \phi]]^y$$




- And the Negative Log Likelihood is:

$$L[\phi] = \sum_{i=1}^I -(1 - y_i) \log[1 - \text{sig}[f[\mathbf{x}_i, \phi]]] - y_i \log[\text{sig}[f[\mathbf{x}_i, \phi]]]$$

- A.k.a. **Binary Cross Entropy Loss**



## Example 2: Binary Classification



# Example 3: Multiclass Classification

- Goal is to assign an input data example  $x$  to one of  $K > 2$  classes, so  $y \in \{1, 2, \dots, K\}$ .
- Real-world examples include: (i) predicting which of  $K = 10$  digits  $y$  is present in an image  $x$  of a handwritten number and (ii) predicting which of  $K$  possible words  $y$  follows an incomplete sentence  $x$ .





# Reminder: Recipe for a Loss

1. Choose a suitable probability distribution  $\Pr(y|\theta)$  defined over the domain of the predictions  $y$  with distribution parameters  $\theta$ .
2. Set the machine learning model  $f[x, \phi]$  to predict one or more of these parameters, so  $\theta = f[x, \phi]$  and  $\Pr(y|\theta) = \Pr(y|f[x, \phi])$ .
3. To train the model, find the network parameters  $\hat{\phi}$  that minimize the negative log-likelihood loss function over the training dataset pairs  $\{x_i, y_i\}$ :

$$\hat{\phi} = \operatorname{argmin}_{\phi} [L[\phi]] = \operatorname{argmin}_{\phi} \left[ - \sum_{i=1}^I \log [Pr(y_i | f[x_i, \phi])] \right]$$


4. To perform inference for a new test example  $x$ , return either the full distribution  $\Pr(y|f[x, \hat{\phi}])$  or the maximum of this distribution.



# Example 3: Multiclass Classification

- We first choose a distribution over the prediction space  $y$ .
- In this case, we have  $y \in \{1, 2, \dots, K\}$ , so we choose the categorical distribution, which is defined on this domain. This has  $K$  parameters  $\lambda_1, \lambda_2, \dots, \lambda_K$ , which determine the probability of each category:  $Pr(y = k) = \lambda_k$ .

- The parameters are constrained to take values between zero and one, and they must collectively sum to one to ensure a valid probability distribution.



# Example 3: Multiclass Classification

- Then we use a network  $\mathbf{f}[\mathbf{x}, \boldsymbol{\phi}]$  with  $K$  outputs to compute these  $K$  parameters from the input  $\mathbf{x}$ .
- Unfortunately, the network outputs will not necessarily obey the sum-to-1 constraints.
  - Consequently, we pass the  $K$  outputs of the network through a function that ensures these constraints are respected.
- A suitable choice is the softmax function. This takes an arbitrary vector of length  $K$  and returns a vector of the same length but where the elements are now in the range  $[0, 1]$  and sum to one. The  $k^{\text{th}}$  output of the softmax function is:

$$\text{softmax}_k[\mathbf{z}] = \frac{\exp[z_k]}{\sum_{k'=1}^K \exp[z_{k'}]}$$



$$\text{softmax}_k[\mathbf{z}] = \frac{\exp[z_k]}{\sum_{k'=1}^K \exp[z_{k']}]$$

## Example 3: Multiclass Classification

- Therefore the likelihood is given by:

$$Pr(y = k|\mathbf{x}) = \text{softmax}_k[\mathbf{f}[\mathbf{x}, \phi]]$$

- And the NLL is given by:

$$\begin{aligned} L[\phi] &= - \sum_{i=1}^I \log \left[ \text{softmax}_{y_i} [\mathbf{f}[\mathbf{x}_i, \phi]] \right] \\ &= - \sum_{i=1}^I \left( f_{y_i} [\mathbf{x}_i, \phi] - \log \left[ \sum_{k'=1}^K \exp [f_{k'} [\mathbf{x}_i, \phi]] \right] \right) \end{aligned}$$



# Colab time: Multiclass cross-entropy loss

<https://colab.research.google.com/drive/1eP81zYuaSWQitBYrNUjnvXc9Yqf6NAjb>



# Multiple Output

- Often, we wish to make more than one prediction with the same model, so the target output  $\mathbf{y}$  is a vector.
  - For example, we might want to predict a molecule's melting and boiling point (a multivariate regression problem, or the object class at every point in an image (a multivariate classification problem).
- While it is possible to define multivariate probability distributions and use a neural network to model their parameters as a function of the input, it is more usual to treat each prediction as independent.

$$Pr(\mathbf{y}|\mathbf{f}[\mathbf{x}_i, \phi]) = \prod_d Pr(y_d|\mathbf{f}_d[\mathbf{x}_i, \phi])$$



# Multiple Output

- The minimization is then done in the usual way by minimizing the Negative Log Likelihood

$$L[\phi] = - \sum_{i=1}^I \log \left[ Pr(\mathbf{y} | \mathbf{f}[\mathbf{x}_i, \phi]) \right] = - \sum_{i=1}^I \sum_d \log \left[ Pr(y_{id} | \mathbf{f}_d[\mathbf{x}_i, \phi]) \right]$$

where  $y_{id}$  is the  $d^{\text{th}}$  output from the  $i^{\text{th}}$  training example.



# Distributions for Loss Functions for Different Prediction Types

Data Type	Domain	Distribution	Use
univariate, continuous, unbounded	$y \in \mathbb{R}$	univariate normal	regression
univariate, continuous, unbounded	$y \in \mathbb{R}$	Laplace or t-distribution	robust regression
univariate, continuous, unbounded	$y \in \mathbb{R}$	mixture of Gaussians	multimodal regression
univariate, continuous, bounded below	$y \in \mathbb{R}^+$	exponential or gamma	predicting magnitude
univariate, continuous, bounded	$y \in [0, 1]$	beta	predicting proportions
multivariate, continuous, unbounded	$\mathbf{y} \in \mathbb{R}^K$	multivariate normal	multivariate regression
univariate, continuous, circular	$y \in (-\pi, \pi]$	von Mises	predicting direction
univariate, discrete, binary	$y \in \{0, 1\}$	Bernoulli	binary classification
univariate, discrete, bounded	$y \in \{1, 2, \dots, K\}$	categorical	multiclass classification
univariate, discrete, bounded below	$y \in [0, 1, 2, 3, \dots]$	Poisson	predicting event counts
multivariate, discrete, permutation	$\mathbf{y} \in \text{Perm}[1, 2, \dots, K]$	Plackett-Luce	ranking





# Why Cross Entropy?

- We have often use the term Cross Entropy to describe the loss commonly used in classification problems, but... Why is it a form of “Entropy”?
- Consider the Kullback-Leibler divergence between two distributions  $p$ , and  $q$ . It measures the “distance” between the two distributions:

$$D_{KL}[q||p] = \int_{-\infty}^{\infty} q(z) \log[q(z)] dz - \int_{-\infty}^{\infty} q(z) \log[p(z)] dz$$

- Now consider that we observe an empirical data distribution at points  $\{y_i\}$ .
  - We can describe this as a weighted sum of point masses:

Dirac's Delta.

$$\delta(x - \alpha) = \frac{1}{2\pi} \int_{-\infty}^{\infty} dp \cos(px - p\alpha)$$

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$



# Why Cross Entropy?

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$

- We want to minimize the KL divergence between  $q(y)$  and the real distribution  $\text{Pr}(y|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \underset{\theta}{\operatorname{argmin}} \left[ \int_{-\infty}^{\infty} q(y) \log[q(y)] dy - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right]\end{aligned}$$



# Why Cross Entropy?

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$

- We want to minimize the KL divergence between  $q(y)$  and the real distribution  $\text{Pr}(y|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \underset{\theta}{\operatorname{argmin}} \left[ \int_{-\infty}^{\infty} q(y) \log[q(y)] dy - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} \left( \frac{1}{I} \sum_{i=1}^I \delta[y - y_i] \right) \log[\text{Pr}(y|\theta)] dy \right]\end{aligned}$$



# Why Cross Entropy?

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$

- We want to minimize the KL divergence between  $q(y)$  and the real distribution  $\text{Pr}(y|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \underset{\theta}{\operatorname{argmin}} \left[ \int_{-\infty}^{\infty} q(y) \log[q(y)] dy - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} \left( \frac{1}{I} \sum_{i=1}^I \delta[y - y_i] \right) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \frac{1}{I} \sum_{i=1}^I \log[\text{Pr}(y_i|\theta)] \right]\end{aligned}$$



# Why Cross Entropy?

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$

- We want to minimize the KL divergence between  $q(y)$  and the real distribution  $\text{Pr}(y|\theta)$ :

$$\begin{aligned}\hat{\theta} &= \underset{\theta}{\operatorname{argmin}} \left[ \int_{-\infty}^{\infty} q(y) \log[q(y)] dy - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} q(y) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \int_{-\infty}^{\infty} \left( \frac{1}{I} \sum_{i=1}^I \delta[y - y_i] \right) \log[\text{Pr}(y|\theta)] dy \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \frac{1}{I} \sum_{i=1}^I \log[\text{Pr}(y_i|\theta)] \right] \\ &= \underset{\theta}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log[\text{Pr}(y_i|\theta)] \right]\end{aligned}$$



# Why Cross Entropy?

$$q(y) = \frac{1}{I} \sum_{i=1}^I \delta[y - y_i]$$

- We want to minimize the KL divergence between  $q(y)$  and the real distribution  $\Pr(y|\theta)$ :

$$\hat{\theta} = \underset{\theta}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log [\Pr(y_i | \theta)] \right]$$

- By substituting  $\theta$  with  $f[\mathbf{x}_i, \phi]$  computed by the ML model we get

$$= \underset{\phi}{\operatorname{argmin}} \left[ - \sum_{i=1}^I \log [\Pr(y_i | f[\mathbf{x}_i, \phi])] \right]$$

- The **negative log-likelihood criterion** (from maximizing the data likelihood) **and the cross-entropy criterion** (from minimizing the distance between the model and empirical data distributions) **are equivalent**.



# Exercise

**Problem 5.6** Consider building a model to predict the number of pedestrians  $y \in \{0, 1, 2, \dots\}$  that will pass a given point in the city in the next minute, based on data  $\mathbf{x}$  that contains information about the time of day, the longitude and latitude, and the type of neighborhood. A suitable distribution for modeling counts is the Poisson distribution (figure 5.15). This has a single parameter  $\lambda > 0$  called the *rate* that represents the mean of the distribution. The distribution has probability density function:

$$Pr(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}. \quad (5.36)$$

Design a loss function for this model assuming we have access to  $I$  training pairs  $\{\mathbf{x}_i, y_i\}$ .



# Solution

- Given the probability density function of the Poisson distribution:  $P_r(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}$





# Solution

- Given the probability density function of the Poisson distribution:  $P_r(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- The log-likelihood for a single data point  $(x, y)$  is:  $\log P_r(y = k) = k \log \lambda - \lambda - \log(k!)$



# Solution

- Given the probability density function of the Poisson distribution:  $P_r(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- The log-likelihood for a single data point  $(x, y)$  is:  $\log P_r(y = k) = k \log \lambda - \lambda - \log(k!)$
- For  $I$  training pairs  $\{x_i, y_i\}$ , the total log-likelihood is the sum of the log-likelihoods for each data point:

$$L(\lambda) = \sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$$



# Solution

- Given the probability density function of the Poisson distribution:  $P_r(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- The log-likelihood for a single data point  $(x, y)$  is:  $\log P_r(y = k) = k \log \lambda - \lambda - \log(k!)$
- For  $I$  training pairs  $\{x_i, y_i\}$ , the total log-likelihood is the sum of the log-likelihoods for each data point:

$$L(\lambda) = \sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$$

- To turn this into a loss function, we take the negative of the log-likelihood:
  - $\text{Loss}(\lambda) = -L(\lambda)$
  - $\text{Loss}(\lambda) = -\sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$



# Solution

- Given the probability density function of the Poisson distribution:  $P_r(y = k) = \frac{\lambda^k e^{-\lambda}}{k!}$
- The log-likelihood for a single data point  $(x, y)$  is:  $\log P_r(y = k) = k \log \lambda - \lambda - \log(k!)$
- For  $I$  training pairs  $\{x_i, y_i\}$ , the total log-likelihood is the sum of the log-likelihoods for each data point:

$$L(\lambda) = \sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$$

- To turn this into a loss function, we take the negative of the log-likelihood:
  - $\text{Loss}(\lambda) = -L(\lambda)$
  - $\text{Loss}(\lambda) = -\sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$
- In practice, the model will predict a different rate  $\lambda_i$  for each data point  $x_i$  based on the input features (time of day, longitude, latitude, type of neighborhood). So, the loss function becomes:

$$\text{Loss}(\lambda) = -\sum_{i=1}^I (y_i \log \lambda - \lambda - \log(y_i!))$$

- The term  $\log(y_i!)$  can be ignored during optimization since it doesn't depend on the model's parameters. It's a constant with respect to  $\lambda$



# Deep Learning

End of Lecture

03 - Model Training. Loss function



SAPIENZA  
UNIVERSITÀ DI ROMA

Fabrizio Silvestri